



Accuracy: every pairwise history test

Read row versus column. Both directions are displayed; each history pair was tested once.


Qwen App_{fr} | 0/120 pairs




Qwen App_{od} | 14/120 pairs




InternVL App_{fr} | 53/120 pairs




InternVL App_{od} | 26/120 pairs



LLaVA App_{fr} | 4/120 pairs



LLaVA App_{od} | 1/120 pairs



Holm-adjusted paired Wilcoxon tests, $p < 0.05$. 120 possible history pairs per tested configuration; 25 paired bags.

■ + Row is significantly better ■ - Row is significantly worse ■ Blank: not significant