



Macro-F1: every pairwise history test

Read row versus column. Both directions are displayed; each history pair was tested once.


Qwen App_{fr} | post-hoc not performed




Qwen App_{od} | post-hoc not performed




InternVL App_{fr} | 33/120 pairs




InternVL App_{od} | 20/120 pairs



LLaVA App_{fr} | post-hoc not performed



LLaVA App_{od} | post-hoc not performed



Holm-adjusted paired Wilcoxon tests, $p < 0.05$. 120 possible history pairs per tested configuration; 25 paired bags.

+ Row is significantly better - Row is significantly worse Blank: not significant